

Zaalima Development - Data Science & AI Division: Q4 Production AI Project Assignments

To: Machine Learning Engineering Track (Cohort Sigma)

From: Principal Data Scientist & AI Architect

Date: November 30, 2025

Subject: Q4 Production AI Project Assignments: Building Production-Grade MLOps Systems----
-A Call to Production

Team, let this be explicitly clear: **Stop training models on the Titanic dataset.** That era is over. The goal of this program is not to produce highly-paid hobbyists. In the real world, data is messy, models drift in production, resource latency matters, and business-critical decisions rely on your code. You are here to build **Production-Grade AI Systems**, not just Jupyter Notebook experiments.

The next four weeks are your final, intensive crucible. We are adopting a strict, non-negotiable **MLOps-first mindset**: if you can't containerize it, monitor it, deploy it, and automatically re-train it, it doesn't exist. Your success will be measured by system resilience and latency, not just the F1-score on a static test set.

The following four projects are designed to cover the entire spectrum of modern Data Science and Machine Learning engineering:

1. **Tabular ML & Time-Series:** Predictive Maintenance (Industrial IoT).
2. **Deep Learning (NLP):** Intelligent Document Processing.
3. **Deep Learning (Computer Vision):** Automated Visual Quality Control.
4. **Big Data Processing:** Scalable Churn Prediction (Apache Spark).

----Contents - Project Index

1. **Tech Stack - Production AI Architecture (The Toolchain)**
2. **Project 1: Manufacturing - Predictive Maintenance System (Time-Series)**
3. **Project 2: Fintech - Intelligent Document Processing (NLP)**
4. **Project 3: Retail - Automated Visual Quality Control (Computer Vision)**
5. **Project 4: Telecom - Big Data Churn Prediction (Apache Spark)**

----1. Tech Stack - The Zaalima Production AI Architecture

We utilize a modern, hybrid stack that surgically balances rapid, local experimentation with scalable, distributed deployment. Mastery of these tools is non-negotiable.

Category	Primary Tools	Purpose & Production Context
Data Foundation	Pandas & NumPy	Vectorization & Local ETL. Used for smaller datasets, high-performance feature vectorization, and data manipulation up to a few GB. Efficient serialization via <code>joblib</code> or <code>pickle</code> is mandatory for moving features to the model.
	Apache Spark (PySpark)	Big Data Processing & Distributed ETL. Essential for Project 4. When data volumes exceed RAM (Terabytes), Spark handles distributed processing, data aggregation, and massive joins across a cluster.
Modeling Engines	Scikit-Learn	Baseline & Feature Pipelines. The robust industry standard for classical ML (Logistic Regression, SVM, Random Forest). Used for establishing initial baselines and building reusable pre-processing pipelines.
	XGBoost/LightGBM	High-Performance Tabular ML. The "Kaggle Winners." Use these gradient boosting frameworks for state-of-the-art results on structured, tabular datasets.
	TensorFlow/Keras	Deep Learning Framework. The workhorse for all complex Neural Network architectures in NLP (LSTMs,

		Transformers) and Computer Vision (CNNs, ResNets).
Development & MLOps	Jupyter Notebooks	Strictly for EDA & Prototyping. Must be refactored into production-ready .py scripts before deployment. A notebook is an exploratory tool, not a deployable artifact.
	PyCharm/VS Code	Production Code IDEs. Used for writing modular, tested, production-ready Python modules and functions.
	MLflow (Bonus)	Experiment & Model Versioning. Highly recommended for tracking all experiments, logging parameters, metrics, and model artifacts for seamless reproducibility.
	Docker	Containerization. Mandatory for Project 2 and Project 4 (simulating a cluster). All API endpoints must be containerized to ensure environment parity between development and production.

----2. Project 1: Manufacturing - IoT Predictive Maintenance Engine

Project Title: IoT Predictive Maintenance Engine (Time-Series Classification)

Product Brand Name: "FactoryGuard AI"

Use Case (Production): A critical manufacturing plant floor contains 500 robotic arms with vibration, temperature, and pressure sensors. The objective is to predict a catastrophic failure **24 hours before it occurs** to allow for scheduled, preemptive maintenance, avoiding millions in unscheduled downtime.

Area	Production Requirements & Features	Implementation Details
Feature Engineering	Advanced Rolling Window Statistics: Creating time-series features such as Rolling Mean, Exponential Moving Average, Standard Deviation of sensor readings over the last 1, 6, and 12 hours. Lag Features ($t-1$, $t-2$) are essential.	Focus on Pandas for feature creation; efficient serialization with <code>joblib</code> .
Modeling	Model Selection: Start with a simple Logistic Regression or Scikit-Learn Random Forest as a baseline. The production model must be XGBoost or LightGBM for maximum predictive power.	Week 2: Hyperparameter Tuning via <code>GridSearchCV</code> or the advanced <code>Optuna</code> library.
Imbalance Handling	Class Imbalance is Extreme: Failures are rare (typically <1% of the data). Do NOT use Accuracy. Must use Precision-Recall Area Under Curve (PR-AUC) . Handle imbalance using SMOTE or, preferably, by adjusting Class Weights in the model.	Use the <code>imbalanced-learn</code> library. Prioritize High Precision (avoiding false alarms).
Explainability	Model Trust & Adoption: Use SHAP (SHapley Additive exPlanations) values to generate local explanations. The maintenance engineer needs to know <i>why</i> the model predicted a failure (e.g., "The rolling mean of temperature	Week 3: Integrate SHAP visualization into your report.


	exceeded 80°C and the variance of vibration tripled").	
Deployment	Real-Time API: Save the final model and pre-processing pipeline using <code>joblib</code> . Build a simple Flask API endpoint that accepts a JSON payload of current sensor readings and returns a failure probability (<code><50ms response time</code>).	Week 4: Latency check is a critical review point.

-----3. Project 2: Fintech - Intelligent Document Processing (NLP)

Project Title: Legal Contract Entity Extractor (NER)

Product Brand Name: "LexiScan Auto"

Use Case (Production): A large financial law firm manages millions of PDF contracts. Manual review is slow and error-prone. The system must automatically extract key, structured entities: **Dates, Party Names, Dollar Amounts, and Termination Clauses** to automatically index and search the documents.



Area	Production Requirements & Features	Implementation Details
Data Acquisition	OCR Integration: The system must handle both native digital and scanned PDF documents . Integrate the Tesseract OCR pipeline to convert pixels to raw, processable text before NLP begins.	Week 1: Focus on text quality and noise reduction from OCR.
Core Modeling	Named Entity Recognition (NER): This is a custom NER task. Training a deep learning model to recognize legal-specific entities.	Use a Bi-Directional LSTM (TensorFlow/Keras) or fine-tune an industrial-grade library like Spacy on your custom annotated data.
Performance & Context	Transfer Learning: Fine-tuning a pre-trained contextual embedding (e.g., GloVe, or a small BERT variant) for superior context awareness, significantly boosting F1-scores.	Week 2: The F1-Score on entity extraction is the key metric. Annotation using tools like Doccano is required for training data.
Post-Processing	Rule-Based Validation: Raw NER output is often imperfect. Implement logic to validate extracted entities (e.g., "Date" must conform to YYYY-MM-DD format; "Amount" must contain currency symbols).	Week 3: Robust handling of edge cases and noise is paramount for lawyer trust.
Deployment	Containerized Microservice: The entire OCR + NLP pipeline (which can be memory-intensive) must be wrapped in a Docker container . Expose the functionality via a resilient REST API endpoint that accepts the PDF/text and	Week 4: The end-to-end test must pass: Upload PDF -> Receive structured JSON.

	returns a JSON payload of extracted entities.	
--	---	--

-----4. Project 3: Retail - Automated Visual Quality Control (Computer Vision)

Project Title: Automated Quality Control (Defect Detection)

Product Brand Name: "VisionSpec QC"

Use Case (Production): A high-speed manufacturing line for PCBs (Printed Circuit Boards) requires real-time quality inspection. The AI must process an image captured by a camera and immediately classify it as "**Pass**" or "**Defect**" (and provide localization of the defect) with high throughput.

Area	Production Requirements & Features	Implementation Details
Data Preparation	<p>Robust Data Augmentation: Real-time data augmentation (rotation, zooming, flipping, brightness shifts) using the Keras <code>ImageDataGenerator</code> during training to ensure the model generalizes across lighting and camera variability.</p>	<p>Week 1: Visualize augmented batches to confirm they are realistic but varied.</p>
Core Modeling	<p>Transfer Learning for Accuracy: Building a CNN from scratch is inefficient. Use a pre-trained, powerful architecture like ResNet50 or a mobile-optimized one like MobileNetV2. Freeze the base layers and fine-tune only the custom classification head.</p>	<p>Week 2: Focus on achieving a low validation loss without overfitting (checked via learning curves).</p>

Interpretability	Visual Verification (Grad-CAM): To build trust, the system must show <i>where</i> it is looking. Implement Grad-CAM (Gradient-weighted Class Activation Mapping) to generate a visual heatmap overlaid on the image, highlighting the specific area of the defect.	Week 3: Verify that heatmaps highlight the actual soldering joint or component, not background noise.
Inference Optimization	High-Speed Inference: The model needs to run at production line speed (e.g., >10 frames per second). Optimize the model artifact (e.g., saving as an optimized <code>.h5</code> or <code>SavedModel</code>).	Week 4: Write a Python script using OpenCV to simulate a live webcam feed and perform frame-by-frame prediction. Live demo is mandatory.

-----5. Project 4: Telecom - Big Data Churn Prediction (Apache Spark)

Project Title: Scalable User Churn Prediction (Big Data ML)

Product Brand Name: "SparkScale Churn"

Use Case (Production): A major telecom company generates **2 Terabytes** of user log data (call duration, data usage, complaints, billing). This volume cannot be loaded into a single machine's RAM. The objective is to build a distributed ML pipeline to predict which high-value users will cancel their subscription next month.

Area	Production Requirements & Features	Implementation Details
Data Processing	Distributed ETL & Storage: Using PySpark RDDs and DataFrames to ingest and process massive, distributed CSV or Parquet files across a simulated cluster environment.	Week 1: Setup a local Spark Cluster (Master + 2 Workers) using Docker . Focus on PySpark schema validation and row counts.

Feature Engineering	Spark SQL for Aggregation: Feature engineering must happen at scale. Utilize Spark SQL to calculate complex user-level aggregates (e.g., "Average data usage in the last 90 days," "Total complaint count"). Use VectorAssembler to combine features.	Week 2: Analyze the Spark Execution Plan (DAG visualization) to identify bottlenecks.
Modeling	Spark MLlib Native: Scikit-Learn cannot be used. We must leverage Spark's native, distributed ML library (pyspark.ml). Train models such as Distributed Logistic Regression and Distributed Random Forest .	Week 3: Setup the BinaryClassificationEvaluator . The focus is on the <i>training time</i> and <i>scalability</i> of the model, not just accuracy.
Deployment	Persisted Pipeline: The entire sequence of transformations (Indexers, Aggregators, Model) must be saved as a single, reusable Spark Pipeline object. This is the production artifact.	Week 4: Write a reusable batch script that loads the saved pipeline, takes in new monthly user data, and generates the final churn probability predictions, simulating a "Daily Batch Job" run.

-----Submission Requirements:

All project assets, including source code, documented pipelines, and the final deployable artifact (Flask API, Docker image, or Spark script), must be available in a dedicated Git repository. The repo link must be submitted to the Zaalima LMS by **Friday 17:00 of Week 4**. Ensure a comprehensive **requirements.txt** file is present to guarantee a fully reproducible environment.

Zaalima Development

Algorithms define reality. Code responsibly.